

# ECE4424: Project Proposal

## A Neural-Network Based Real-Time SISO Controller

Mihir Savadi  
Bradley Department of Electrical and Computer Engineering  
Virginia Tech

mihirsavadil@vt.edu

### 1. Introduction

PID (potential, integral, derivative) control loops are ubiquitous in our world today – it is the de-facto control mechanism behind ovens, toasters, home or auto HVACs, cruise controls, auto pilots, attitude control in drones, and various industrial processes. Its popularity is due to its simplicity, reliability, and interpretability – it has been used in industry for several decades already. One drawback of the PID control mechanism – and this is the focus of my project – is that the parameters which govern them require manual human tuning via trial and error. In many applications this may be either extremely inconvenient, costly, or just infeasible. I aim to design and test an approach employing Neural-Networks (NN) that, for any system thrown under its purview, can automatically converge upon equivalent behavior of an ideally tuned PID control loop, given real-time data fed to it from the system its controlling and no human input.

For those unfamiliar, I will first outline what PID controller's are – they are remarkably simple. I will then explain my motivations for this project, after which I will detail my preliminary approach and evaluation metrics. Finally I will discuss the tasks and milestones I will have to achieve in order to realize my idealized result – an adaptive real-time alternative to PID controllers.

### 2. PID Controllers Explained

Say we had an oven that we wanted to control the temperature of. The intended temperature can be set to any value in a given range by a cook, and the actual temperature of the oven would then be influenced by a heating coil within it. This heating coil would be controlled by some electrical circuitry (often times mechanical!<sup>1</sup>) that would carefully control the power of the heating coil such that a steady temperature is maintained. If the oven is already at

---

<sup>1</sup>Bimetallic strips are commonly found controlling our indicators in our cars – making sure their switching frequency is constant; or the temperature control in mini-fridges.

a steady temperature, and the cook increases the intended temperature, the controller would then be responsible for ramping up power of the coil so that the oven is able to settle at the new temperature quickly, but not overshoot or oscillate around it (i.e. it must ensure a *perfectly damped* response). The inverse of the same logic applies if the cook decreased the temperature instead.

In this example, our oven is our system that we are concerned about, which we will refer to as the **plant**. The output characteristic of our plant that we wish to control is the temperature, which varies with time – we will refer to it as  $y(t)$ . The input to our plant that our controller control's in order to influence  $y(t)$  is the power of the heating coil, which also varies with time – we will refer to it as  $u(t)$ . The intended value of  $y(t)$  – as set by any agent, in our example the cook sets an intended temperature – can also vary with time and will be referred to as  $r(t)$ . The error between the intended and the actual plant output will be referred to as  $e(t)$ . Eq 1 below defines a PID controller.

$$u(t) = K_P e(t) + K_I \int_0^t e(\tau) d\tau + K_D \frac{de(t)}{dt} \quad (1)$$

where  $e(t) = r(t) - y(t)$

In Section 1 we mentioned how PID controllers have parameters that require manual tuning – these are  $K_P$  (called the potential),  $K_I$  (called the integral), and  $K_D$  (called the derivative), which can be seen in Eq 1 above. Figure 1 is a block diagram representation of Eq 1, which should help illustrate the real-time nature of this type of system. It can be observed that PID controllers are Single-Input Single-Output systems (SISO).

Considering that our controller will be digital, we will be operating in discrete time. As such, eq 2 below is the 'digital' version of eq 1. As can be seen in eq 2, this PID controller that we wish to emulate with our NN based approach will require memory (of up to  $n$  samples prior to the most recent sample) of temporal information.




Figure 1. Block diagram of a PID controller. *Source: Arturo Urquiza* - [https://en.wikipedia.org/wiki/PID\\_controller#/media/File:PID\\_en.svg](https://en.wikipedia.org/wiki/PID_controller#/media/File:PID_en.svg)

$$u[t] = K_P e[t] + K_I \sum_{i=0}^n e[t-i] + K_D (e[t] - e[t-1])$$

where  $e[t] = r[t] - y[t]$ ,  
and  $n \geq 0$

(2)

### 3. Motivation

In certain situations, relying on PID control systems to achieve control stability is impractical. A common example would be when one-off systems, that aren't mass produced and don't have R&D time before mass production, need to be tuned. For example, whenever an amateur racing drone pilot builds a new custom drone with an off-the-shelf flight controller, they then assume the burden of tuning the PID parameters that govern their drone's flight characteristics. This involves flying around, observing flight behavior, and tuning iteratively – this is clearly cumbersome and can even lead to situations that may damage this one-off drone. Modern flight control software comes with auto-tuning features, but these are mostly deterministic and perform poorly.

Another example would be DIY reflow ovens. When hobbyists or independent printed circuit board designers want to build large complex boards or boards with several hard-to-hand-solder surface mount devices, they often seek the refuge of 'reflow-ovens'. The reflow oven process involves a PCB with components placed onto their footprints with unsoldered 'solder paste' in between. This PCB is then placed into a reflow oven, which controls the temperature in its chamber to follow a very specific temperature-time curve in order for the solder paste to melt and create sound solder joints as per the manufacturers specifications. Often times hobbyists modify toaster ovens into reflow ovens (which are very popular and work remarkably well) by swapping out their cheap mechanical controls with their own programmed PID controllers, which they then have to tune. If instead they could use an automatic 'adaptive' controller, it would make the conversion process of any randomly chosen toaster oven a lot easier.

My personal motivation is this: I thoroughly enjoy own-

ing and working on old cars. I recently bought a 2004 Land Rover Discovery II, and have since been introduced to a few people involved in the North American Land Rover owners communities. One of these people creates and sells kits to modernize and improve the reliability of certain subsystems of Discovery II's in particular. He has recently been trying to create an independent electrically driven air cooling system that would replace the traditional belt-driven fan systems that ship with not just Discovery II's, but any other car as well. This system would as a result need to monitor engine temperatures and adapt the power of its electrically driven fans and other cooling systems to match the thermal characteristics of whatever car the kit happens to be installed in, ideally without requiring the user to have to sit through a long and painful tuning process. From another perspective, instead of testing and shipping pre-tuned controllers for only a predetermined fixed set of cars – which would require much more manual labor on the creators end – an automatic 'adaptive' controller would not only prevent this manual labor but also expand the potential market for the kit as well.

Similar arguments from the examples above can be used to justify the benefits of such an 'adaptive' controller in various industrial applications. As such, many efforts have been made and are currently even being implemented in deep industry – please see the references section for a list of all the sources I have explored within this space. The approach I will be designing (detailed in Section 4) is not one that I have seen yet in existing literature, and is possibly a lot simpler than many of the approaches I have read so far.

### 4. Approach

The first question that arises is, on what plant<sup>2</sup> do I test the controller I intend to create? Initially I was going to modify a toaster oven's control circuitry with my own PCB involving an SSR (solid state relay), micro-controller, and other supporting electronics (at the time of writing this I already bought a toaster oven in anticipation, and need to return it). However upon further research, I found that engineers quite often emulate plants with computer models in order to facilitate PID and MPC (Model Predictive Controllers) control loops in software before physical testing. One of the models they use is called the First Order Linear System with Time Delay (FOPDT)<sup>3 4</sup>. Many open-source python based FOPDT simulators exist<sup>5</sup>, which I can easily modify to fit my own needs and controller to. This will pro-

<sup>2</sup>refer to Section 2 for what 'plant' means

<sup>3</sup><https://apmonitor.com/pdc/index.php/Main/FirstOrderSystems>

<sup>4</sup><https://apmonitor.com/pdc/index.php/Main/FirstOrderOptimization>

<sup>5</sup>[https://github.com/Destination2Unknown/PythonPID\\_Simulator](https://github.com/Destination2Unknown/PythonPID_Simulator)

vide a much more efficient and reliable platform to test my approach before diving into any hardware testing.

The high-level structure of the NN architecture I envision is as follows. Please keep in mind that the notation used here will mirror what was used in Section 2. It will have an input layer consisting of multiple neurons; a single hidden layer consisting of an unknown number of neurons that I will have to experiment with; and an output layer consisting of a single neuron that will function as  $u[t]$ .

The input layer will consist of neurons, each of which will receive inputs described in eq 3 as follows.

$$\{r[t - k], u[t - i], y[t - i]\} \quad (3)$$

where  $0 \leq k \leq n$  and  $1 \leq i \leq n$

The NN will be initialized with random weights and biases, and will undergo a specified number of rounds of back propagation on its entire data-set in real-time, that is at every time-step  $t$ . There will be no training data when the controller first initializes, but at every time-step  $t$ , a labelled data point will be appended to the training set – it will have ‘inputs’ as described in eq 3, ‘output’  $u[t]$ , and error value (i.e. the objective function that the gradient descent algorithm will attempt to minimize) as described in eq 4 below.

$$e[t] = \frac{1}{z + 1} \sum_{i=0}^z (r[i] - y[i])^2 \quad (4)$$

where  $0 \leq z \leq n$

The  $z$  variable in eq 4 points to a time step that is somewhere in between the most recent time step and the beginning of time. It functions as a metric to optimize for perfect dampening of  $u[t]$  to  $r[t]$  within some specified time window  $z$ . If  $u[t]$  is under damped, it will oscillate around  $r[t]$  before converging, and will take a long time to converge to  $r[t]$ , which is what eq 4 penalizes. Similarly, if  $u[t]$  is over damped, it will approach  $r[t]$  very slowly and also take a long time to converge to  $r[t]$ , which eq 4 also penalizes. Ideally, for a given time window  $z$ ,  $e[t]$  would be as close to zero as possible.

Within some margin for error,  $z$  needs to be adjusted to be compatible with the typical rate of change of  $r[t]$  that the controller would experience in a particular environment. Eq 4 becomes less effective in quantifying quality of dampening if  $r[t]$  varies significantly between time-steps  $t$  and  $t - z$ , as less time is given for the controller to ‘adapt’ to changes in its input – it is unrealistic to expect a plant to ‘snap’ to an input  $u[t]$ ; it would be impossible for an FOPDT modelled plant. Including current and previous samples of  $r[t]$  as inputs to the NN (as described in eq 3) will hopefully help with this.

Over time, as the training set grows, the NN would adapt and ‘converge’ towards close-to-perfectly damped behavior. If changes in the environment of the plant are expected,

then this would be reflected by varying the parameters that govern the FOPDT plant simulation model in time. Training data collected before a pre-determined point in time can then be deleted or ‘forgotten’, in order to allow for the NN controller to adapt.

The ultimate goal would be to create a NN model that converges as quickly as possible to near perfect damping for a variety of FOPDT plant model parameter ranges. The goal would be to reduce the frequency of time-steps as much as possible in order to reduce the NN controller’s latency which would influence its ability to ‘sense’ the plant and respond quickly. Time-step frequency would be limited by how quickly the processor that the NN is running on can do the specified number of rounds of back propagation. A lot of variables can influence compute time and therefore also the time-step frequency – the size of the training data set (which can be improved by controlling when old training data is ‘forgotten’);  $k$  and  $n$  from eq 3 and the number of neurons in the hidden layer directly control the number of weights and biases we have to deal with; the activation function we use (e.g. ReLU is much less computationally intensive than sigmoid); how large  $z$  is from eq 4. On the flip side, we could also argue that there could be a point of marginal returns in time-step frequency reduction due to NN complexity being insufficient or training data size being too small. Conversely, if time-step frequency is too high,  $k$  and  $n$  from eq 3 would need to be higher in order to capture a significantly large enough  $z$  is from eq 4, increasing compute. All these nuances will have to be experimented with.

We could further speed up convergence time, instead of starting off with random weights and biases we could train the NN to mimic a poor performing but very general purpose PID controller. When the error value from eq 4 becomes small enough, we can cease training and just do inference, which would free up compute significantly.

## 5. Evaluation Metrics and Milestones

The following is a chronological list of high-level steps needed to be accomplished to achieve the desired goal of this project – a functioning NN network that would converge to the performance of a manually hand tuned PID controller for a variety of types of FOPDT emulated plants.

1. Build FOPDT simulator code base with good abstraction so that a variety of FOPDT and controller varieties can be accommodated.
2. Create a collection of FOPDT simulation scenarios and performance measurement system for automated controller performance testing.
3. Write, test, and bench mark a PID controller method in the simulator from the previous step.

4. Write a general back propagation algorithm, or find an optimized solution off-the-shelf that is compatible with any arbitrary model.
5. Write and test the novel NN controller code within the simulator from test 1. Ensure that all the variables discussed at the end of section 4 are easily adjustable for the sake of experimentation.
6. Compile results and comparisons from steps 3 and 5 and report observations in final report.

## References

- [1] Andrew G. Barto. Reinforcement learning control. *Current Opinion in Neurobiology*, 4(6):888–893, 1994.
- [2] Silvia Ferrari and Robert F. Stengel. *An adaptive critic global controller*, volume 4, pages 2665–2670. 2002.
- [3] Martin T. Hagan, Howard B. Demuth, and Orlando De Jesús. An introduction to the use of neural networks in control systems. *International Journal of Robust and Nonlinear Control*, 12(11):959–985, 2002.
- [4] Andrej Karpathy. The unreasonable effectiveness of recurrent neural networks. <https://karpathy.github.io/2015/05/21/rnn-effectiveness/>, 2015.
- [5] F.W. Lewis, S. Jagannathan, and A. Yesildirak. *Neural Network Control Of Robot Manipulators And Non-Linear Systems*. CRC Press, 2020.
- [6] Nicholas Lewis. Emulating a pid controller with long short-term memory, series. [towardsdatascience.com](https://towardsdatascience.com), 2020.
- [7] Derong Liu. An introduction to adaptive critic control: A paradigm based on approximate dynamic programming. 2006.
- [8] W. Thomas Miller, Richard S. Sutton, and Paul J. Werbos. *Computational Schemes and Neural Network Models for Formation and Control of Multijoint Arm Trajectory*, pages 197–228. 1995.
- [9] D.H. Nguyen and B. Widrow. Neural networks for self-learning control systems. *IEEE Control Systems Magazine*, 10(3):18–23, 1990.
- [10] Christopher Olah. Understanding lstm networks. <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>, 2015.
- [11] Steven Spielberg Pon Kumar, Aditya Tulsyan, Bhushan Gopaluni, and Philip Loewen. A deep learning architecture for predictive control. *IFAC-PapersOnLine*, 51(18):512–517, 2018. 10th IFAC Symposium on Advanced Control of Chemical Processes ADCHEM 2018.
- [12] Jun Zhao, Qingliang Zeng, and Bin Guo. Adaptive critic learning-based robust control of systems with uncertain dynamics. *Computational Intelligence and Neuroscience*, 2021:1–8, 11 2021.

## License

This work is licensed under a Creative Commons “Attribution-NonCommercial-ShareAlike 4.0 International” license.

